

```

1 import os
2 import sys
3 import scipy.io
4 import scipy.misc
5 import matplotlib.pyplot as plt
6 from matplotlib.pyplot import imshow
7 from PIL import Image
8 import numpy as np
9 import tensorflow as tf
10

```

Neural Style Transfer merges two images, namely: a "content" image (C) and a "style" image (S), to create a "generated" image (G). The generated image G combines the "content" of the image C with the "style" of image S.

Neural Style Transfer (NST) uses a previously trained convolutional network, and builds on top of that. The idea of using a network trained on a different task and applying it to a new task is called transfer learning. Used VGG-19, a 19-layer version of the VGG network. This model has already been trained on the very large ImageNet database, and has learned to recognize a variety of low level features (at the shallower layers) and high level features (at the deeper layers).

```

1 vgg = tf.keras.applications.VGG19(include_top=False,
2                                     input_shape=(img_size, img_size, 3),
3                                     weights='imagenet')
4
5 vgg.trainable = False
<keras.src.engine.functional.Functional object at 0x7a111f397f10>

```

Building the Neural Style Transfer (NST) algorithm in three steps:

- 1) content cost function  $J_{content}(C, G)$
- 2) style cost function  $J_{style}(S, G)$
- 3) Total Cost  $J(G) = \alpha * J_{content}(C, G) + \beta * J_{style}(S, G)$

$$J_{content}(C, G) = \frac{1}{4 \times n_H \times n_W \times n_C} \sum_{\text{all entries}} (a^{(C)} - a^{(G)})^2$$

$$J_{style}^{[l]}(S, G) = \frac{1}{4 \times n_C^2 \times (n_H \times n_W)^2} \sum_{i=1}^{n_C} \sum_{j=1}^{n_C} (G_{(gram)i,j}^{(S)} - G_{(gram)i,j}^{(G)})^2$$

- The shallower layers of a ConvNet tend to detect lower-level features such as edges and simple textures.
- The deeper layers tend to detect higher-level features such as more complex textures and object classes.

We need the "generated" image G to have similar content as the input image C. Suppose you have chosen some layer's activations to represent the content of an image.

- will get the most visually pleasing results if you choose a layer from somewhere in the middle of the network—neither too shallow nor too deep. This ensures that the network detects both higher-level and lower-level features.

forward propagate image C,

- Set the image C as the input to the pretrained VGG network, and run forward propagation.
- Let  $a^{(C)}(l)$  be the hidden layer activations in the layer l. This will be an  $n_H \times n_W \times n_C$  tensor.

forward propagate image G,

- Repeat this process with the image G: Set G as the input, and run forward propagation.
- Let  $a^G(l)$  be the corresponding hidden layer activation.

```
1 content_image = Image.open("/content/baby.jpg")
2 content_image
3
```



```

1 #function to compute Content Cost
2 def compute_content_cost(content_output, generated_output):
3     """
4     Computes the content cost
5
6     Arguments:
7     a_C -- tensor of dimension (1, n_H, n_W, n_C), hidden layer activations representing content of the image C
8     a_G -- tensor of dimension (1, n_H, n_W, n_C), hidden layer activations representing content of the image G
9
10    Returns:
11    J_content -- scalar
12    """
13    a_C = content_output[-1]
14    a_G = generated_output[-1]
15
16    m, n_H, n_W, n_C = a_G.get_shape().as_list()
17
18    # Reshape 'a_C' and 'a_G'
19    a_C_unrolled = tf.reshape(a_C, shape=[m, n_H * n_W, n_C])
20    a_G_unrolled = tf.reshape(a_G, shape=[m, n_H * n_W, n_C])
21
22    J_content = tf.reduce_sum(tf.square(a_C_unrolled - a_G_unrolled))/(4.0 * n_H * n_W * n_C)
23
24    return J_content

```

```

1 example = Image.open("/content/Artwork.jpg")
2 example
3

```



#### Style Matrix/Gram matrix

- For the sake of clarity, in this assignment G\_gram will be used to refer to the Gram matrix, and G to denote the generated image.

Computing Gram matrix G\_gram by multiplying the "unrolled" filter matrix with its transpose:

```
G_gram = A_unrolled * A_unrolled^T
```

```

1 #function to Compute Gram Matrix
2 def gram_matrix(A):
3     """
4     Argument:
5     A -- matrix of shape (n_C, n_H*n_W)
6     Returns:
7     GA -- Gram matrix of A, of shape (n_C, n_C)
8     """
9     GA = tf.matmul(A, tf.transpose(A))
10    return GA

```

```

1 #function to compute layer style cost
2 def compute_layer_style_cost(a_S, a_G):
3     """
4     Arguments:
5     a_S -- tensor of dimension (1, n_H, n_W, n_C), hidden layer activations representing style of the image S
6     a_G -- tensor of dimension (1, n_H, n_W, n_C), hidden layer activations representing style of the image G
7
8     Returns:
9     J_style_layer
10    """
11    m, n_H, n_W, n_C = a_G.get_shape().as_list()
12
13    # Reshape the tensors from (1, n_H, n_W, n_C) to (n_C, n_H * n_W)
14    a_S = tf.transpose(tf.reshape(a_S, shape=[-1, n_C]))
15    a_G = tf.transpose(tf.reshape(a_G, shape=[-1, n_C]))
16
17    # Computing gram_matrices for both images S and G
18    GS = gram_matrix(a_S)
19    GG = gram_matrix(a_G)
20
21    J_style_layer = tf.reduce_sum(tf.square(GS - GG))/(4.0 *((n_H * n_W * n_C)**2))
22
23
24    return J_style_layer

```

### Style Weights

- better results if style costs from several different layers are merged.
- Each layer will be given weights, that reflect how much each layer will contribute to the style.

$$J_{style}(S, G) = \sum_l \lambda^{[l]} J_{style}^{[l]}(S, G)$$

listing the layer names:

```

1 for layer in vgg.layers:
2     print(layer.name)

input_4
block1_conv1
block1_conv2
block1_pool
block2_conv1
block2_conv2
block2_pool
block3_conv1
block3_conv2
block3_conv3
block3_conv4
block3_pool
block4_conv1
block4_conv2
block4_conv3
block4_conv4
block4_pool
block5_conv1
block5_conv2
block5_conv3
block5_conv4
block5_pool

```

output of layer 'block5\_conv4'.the content layer, which will represent the image.

```
1 vgg.get_layer('block5_conv4').output
```

```
<KerasTensor: shape=(None, 25, 25, 512) dtype=float32 (created by layer 'block5_conv4')>
```

layers to represent the style of the image and assigning style costs:

```

1 STYLE_LAYERS = [
2     ('block1_conv1', 0.2),
3     ('block2_conv1', 0.2),
4     ('block3_conv1', 0.2),
5     ('block4_conv1', 0.2),
6     ('block5_conv1', 0.2)]
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
```

```

1 #function to compute combined style cost
2
3 def compute_style_cost(style_image_output, generated_image_output, STYLE_LAYERS=STYLE_LAYERS):
4     """
5         Computes the overall style cost from several layers
6
7         Arguments:
8             style_image_output -- our tensorflow model
9             generated_image_output --
10            STYLE_LAYERS -- A python list containing:
11                - the names of the layers we would like to extract style from
12                - a coefficient for each of them
13
14        Returns:
15            J_style
16            """
17
18
19    J_style = 0
20
21    # Set a_S to be the hidden layer activation from the layer we have selected.
22    # The last element of the array contains the content layer image
23    a_S = style_image_output[:-1]
24
25    # Set a_G to be the output of the hidden layers.
26    # The last element of the list contains the content layer image
27    a_G = generated_image_output[:-1]
28    for i, weight in zip(range(len(a_S)), STYLE_LAYERS):
29        # Compute style_cost for the current layer
30        J_style_layer = compute_layer_style_cost(a_S[i], a_G[i])
31
32        # Add weight * J_style_layer of this layer to overall style cost
33        J_style += weight[1] * J_style_layer
34
35    return J_style
```

Total\_cost

$$J(G) = \alpha J_{content}(C, G) + \beta J_{style}(S, G)$$

```

1 #function to Compute Total Cost
2 @tf.function()
3 def total_cost(J_content, J_style, alpha = 10, beta = 40):
4     """
5     Arguments:
6     J_content -- content cost coded above
7     J_style -- style cost coded above
8     alpha -- hyperparameter weighting the importance of the content cost
9     beta -- hyperparameter weighting the importance of the style cost
10
11    Returns:
12    J -- total cost
13    """
14    J = alpha * J_content + beta * J_style
15
16    return J

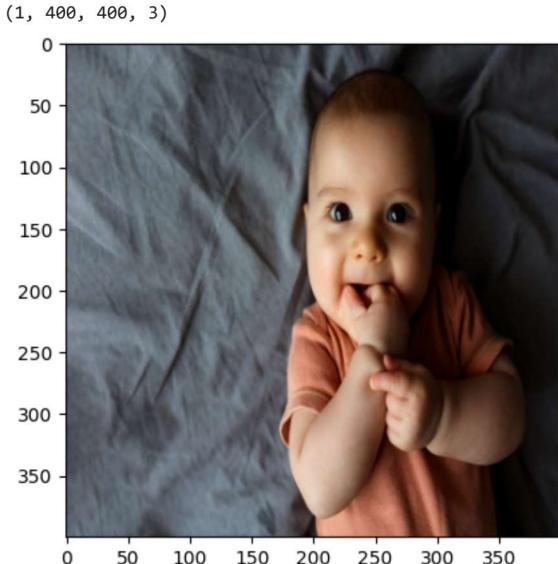
```

Load the Content Image and Resizing it to (400,400) and converting it to a tensor

```

1 content_image = np.array(Image.open("/content/baby.jpg").resize((400,400)))
2 content_image = tf.constant(np.reshape(content_image, ((1,) + content_image.shape)))
3
4 print(content_image.shape)
5 imshow(content_image[0])
6 plt.show()

```



Load the Style Image and resizing and converting it to a tensor

```

1 style_image = np.array(Image.open("/content/Artwork.jpg").resize((img_size, img_size)))
2 style_image = tf.constant(np.reshape(style_image, ((1,) + style_image.shape)))
3
4 print(style_image.shape)
5 imshow(style_image[0])
6 plt.show()

```

```
(1, 400, 400, 3)
```



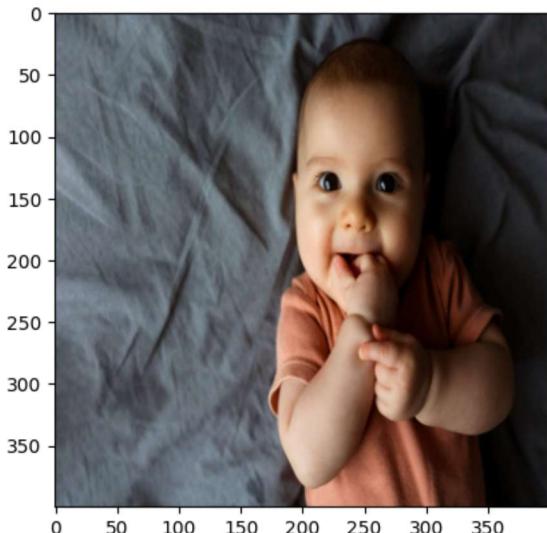
Randomly initialize the image to be Generated

- initializing the "generated" image as content\_image.



```
1 generated_image = tf.Variable(tf.image.convert_image_dtype(content_image, tf.float32))
2 # noise = tf.random.uniform(tf.shape(generated_image), -0.25, 0.25)
3 # generated_image = tf.add(generated_image, noise)
4 # generated_image = tf.clip_by_value(generated_image, clip_value_min=0.0, clip_value_max=1.0)
5
6 print(type(generated_image))
7 imshow(generated_image.numpy()[0])
8 plt.show()
```

```
<class 'tensorflow.python.ops.resource_variable_ops.ResourceVariable'>
```



Loading Pre-trained VGG19 Model

```
1 def get_layer_outputs(vgg, layer_names):
2     """ Creates a vgg model that returns a list of intermediate output values."""
3     outputs = [vgg.get_layer(layer[0]).output for layer in layer_names]
4
5     model = tf.keras.Model([vgg.input], outputs)
6     return model
```

defining the content layer and build the model.

```
1 content_layer = [('block5_conv4', 1)]
2
3 vgg_model_outputs = get_layer_outputs(vgg, STYLE_LAYERS + content_layer)

1 content_target = vgg_model_outputs(content_image) # Content encoder
2 style_targets = vgg_model_outputs(style_image) # Style encoder
```

Compute the Content image Encoding (a\_C)

```

1 # Assign the content image to be the input of the VGG model.
2 # Set a_C to be the hidden layer activation from the layer we have selected
3 preprocessed_content = tf.Variable(tf.image.convert_image_dtype(content_image, tf.float32))
4 a_C = vgg_model_outputs(preprocessed_content)

```

1

Compute the Style image Encoding (a\_S)

The code below sets a\_S to be the tensor giving the hidden layer activation for 'STYLE\_LAYERS' using our style image.

```

1 # Assign the input of the model to be the "style" image
2 preprocessed_style = tf.Variable(tf.image.convert_image_dtype(style_image, tf.float32))
3 a_S = vgg_model_outputs(preprocessed_style)

1 def clip_0_1(image):
2     """
3         Truncate all the pixels in the tensor to be between 0 and 1
4     """
5     return tf.clip_by_value(image, clip_value_min=0.0, clip_value_max=1.0)

1 def tensor_to_image(tensor):
2     """
3         Converts the given tensor into a PIL image
4     """
5     tensor = tensor * 255
6     tensor = np.array(tensor, dtype=np.uint8)
7     if np.ndim(tensor) > 3:
8         assert tensor.shape[0] == 1
9         tensor = tensor[0]
10    return Image.fromarray(tensor)

```

Solving Optimization Problem

Implementing the train\_step() function for transfer learning

- Using the Adam optimizer to minimize the total cost J.
- Using a learning rate of 0.01
- tf.GradientTape to update the image.
- tf.GradientTape():
  - Computing the encoding of the generated image using vgg\_model\_outputs. Assign the result to a\_G.
  - Computing the total cost J, using the global variables a\_C, a\_S and the local a\_G
  - Using 'alpha = 10' and 'beta = 40'.

```

1 print("TensorFlow version:", tf.__version__)
TensorFlow version: 2.14.0

1 optimizer = tf.keras.optimizers.Adam(learning_rate=0.001)
2 @tf.function()
3 def train_step(generated_image, alpha = 10, beta = 40):
4     # generated_image = tf.Variable(initial_value=generated_image, trainable=True)
5     with tf.GradientTape() as tape:
6         tape.watch(generated_image)
7         # Computing a_G as the vgg_model_outputs for the current generated image
8         a_G = vgg_model_outputs(generated_image)
9         J_style = compute_style_cost(a_S, a_G)
10        # print(J_style.shape)
11        J_content = compute_content_cost(a_C, a_G)
12        # print(J_content.shape)
13        # Computing the total cost
14        J = total_cost(J_content, J_style, alpha = alpha, beta = beta)
15        print(type(J), J.shape)
16        # print(J.shape)
17
18
```

```

```
19     grad = tape.gradient(J, generated_image)
20     print(type(grad),grad.shape)
21     print(type(generated_image),generated_image.shape)
22     optimizer.apply_gradients([(grad, generated_image)])
23     generated_image.assign(clip_0_1(generated_image))
24 # return J
```

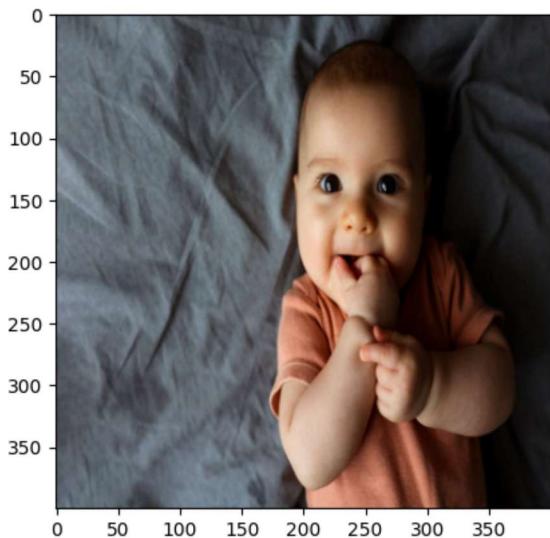
Update the image in every epoch- Training Step

```
1 # Showing the generated image at some epochs
2 epochs = 20001
3 for i in range(epochs):
4     train_step(generated_image)
5     if i % 500 == 0:
6         print(f"Epoch {i} ")
7         image = tensor_to_image(generated_image)
8         imshow(image)
9         plt.show()
```

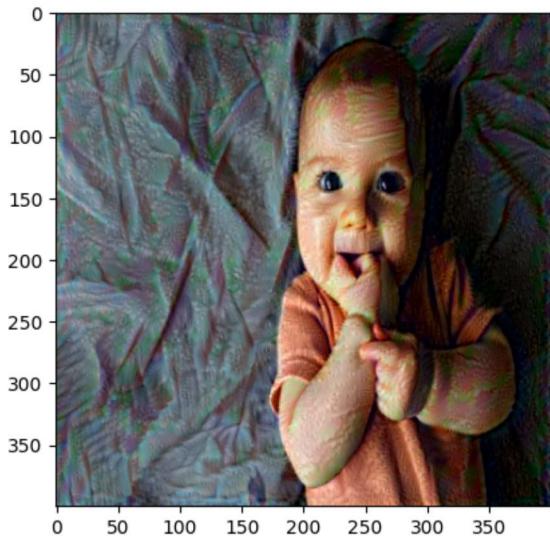


```
<class 'tensorflow.python.framework.ops.Tensor'> ()  
<class 'tensorflow.python.framework.ops.Tensor'> (1, 400, 400, 3)  
<class 'tensorflow.python.ops.resource_variable_ops.ResourceVariable'> (1, 400, 400, 3)  
<class 'tensorflow.python.framework.ops.Tensor'> ()  
<class 'tensorflow.python.framework.ops.Tensor'> (1, 400, 400, 3)  
<class 'tensorflow.python.ops.resource_variable_ops.ResourceVariable'> (1, 400, 400, 3)
```

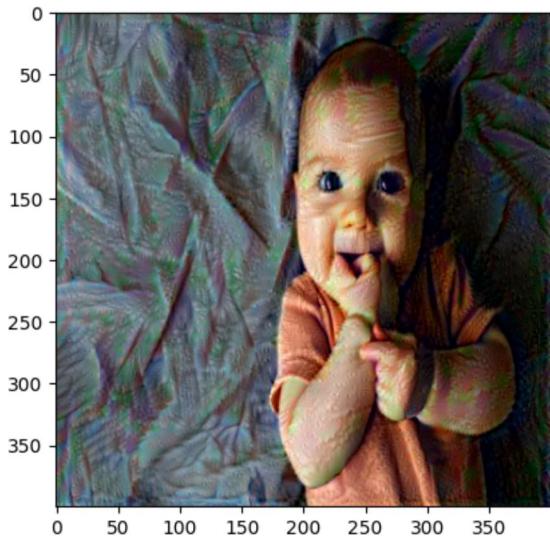
Epoch 0



Epoch 500



Epoch 1000



Epoch 1500

